

PFMIS Bundled Local AI Setup Guide

Use this after the large downloads finish.

This guide explains how to place `llama-server.exe` and the `.gguf` model where PFMIS expects them.

PFMIS project folder:

```
D:\School\JAVA\MyFin_MIS
```

1. Files You Need

- [] llama.cpp Windows files, including `llama-server.exe`
- [] One GGUF model file, for example `Llama-3.2-3B-Instruct-Q4_K_M.gguf`
- [] Licence files for llama.cpp and the model

2. Copy llama-server.exe

Open the downloaded llama.cpp ZIP file. Copy `llama-server.exe` into this folder:

```
D:\School\JAVA\MyFin_MIS\local-ai\runtime\
```

If the llama.cpp ZIP contains DLL files, copy those DLL files into the same folder as `llama-server.exe`.

Expected result

```
D:\School\JAVA\MyFin_MIS\local-ai\runtime\llama-server.exe
```

3. Copy The Model

Find the downloaded model file. It must end with `.gguf`.

Rename the model file to:

```
pfmis-model.gguf
```

Copy it into this folder:

```
D:\School\JAVA\MyFin_MIS\local-ai\models\
```

Expected result

```
D:\School\JAVA\MyFin_MIS\local-ai\models\pfmis-model.gguf
```

Important: Make sure the final filename is exactly `pfmis-model.gguf`. It should not become `pfmis-model.gguf.gguf` or `pfmis-model.gguf.txt`.

4. Confirm The Folder Structure

After copying, the folders should look like this:

```
D:\School\JAVA\MyFin_MIS
|-- local-ai
    |-- runtime
    |   |-- llama-server.exe
    |   |-- other DLL files, if included
    |-- models
    |   |-- pfmis-model.gguf
    |-- agents
    |   |-- pfmis-finance-agent.json
    |-- licenses
        |-- llama.cpp licence
        |-- model licence
```

5. Start PFMIS

1. Open PFMIS.
2. Go to **Settings**.
3. Open **AI Registration**.
4. Select **Bundled Local AI**.
5. Make sure **Enable AI recommendations** is checked.

6. Click **Save AI**.

7. Click **Test AI**.

PFMIS will automatically use:

```
Provider: Bundled Local AI  
Endpoint: http://127.0.0.1:11435/v1  
Model: pfmis-local  
API key: Not required
```

6. What Happens When You Test AI

PFMIS will start this command internally:

```
local-ai\runtime\llama-server.exe --model local-ai\models\pfmis-model.gguf --alias pfmis-local --host 127.0.0.1 --port 11435 --ctx-size 4096
```

The first start can be slow because the model is large. Wait until the status says the local AI is ready.

7. Use AI In Goals

1. Go to **Goals**.
2. Register or enter a goal.
3. Click **Generate AI Advice**.
4. The AI will fill Achievement Plan, Spending Insight, Where To Adjust, and Advice.

8. Troubleshooting

PFMIS says the runtime is missing

Check that this file exists:

```
D:\School\JAVA\MyFin_MIS\local-ai\runtime\llama-server.exe
```

PFMIS says the model is missing

Check that this file exists:

```
D:\School\JAVA\MyFin_MIS\local-ai\models\pfmis-model.gguf
```

The model takes long to load

This is normal for large files. A smaller quantized model such as Q4 is better for older computers.

The computer becomes slow

Use a smaller model. For example, use a 1B or 3B GGUF model before trying an 8B model.

Antivirus blocks llama-server.exe

Allow the file only if it came from the official llama.cpp release and you trust the download source.

Port 11435 is already in use

Close any old local AI process, then open PFMIS again.

9. Licence Reminder

Before distributing PFMIS with the model, confirm that the selected model licence allows redistribution, offline installation, and your intended use.

The llama.cpp runtime licence and the model licence are separate. Include both licence notices in:

```
D:\School\JAVA\MyFin_MIS\local-ai\licenses\
```

Final Checklist

- [] llama-server.exe is inside local-ai\runtime
- [] Required DLL files are inside local-ai\runtime
- [] The model is renamed to pfmis-model.gguf
- [] The model is inside local-ai\models
- [] PFMIS provider is set to **Bundled Local AI**
- [] **Enable AI recommendations** is checked
- [] **Test AI** passes